



Analysis of ByteTrack and YOLOX on Various Input Sizes

Sara Larson and Ethan Sims

Selected Application

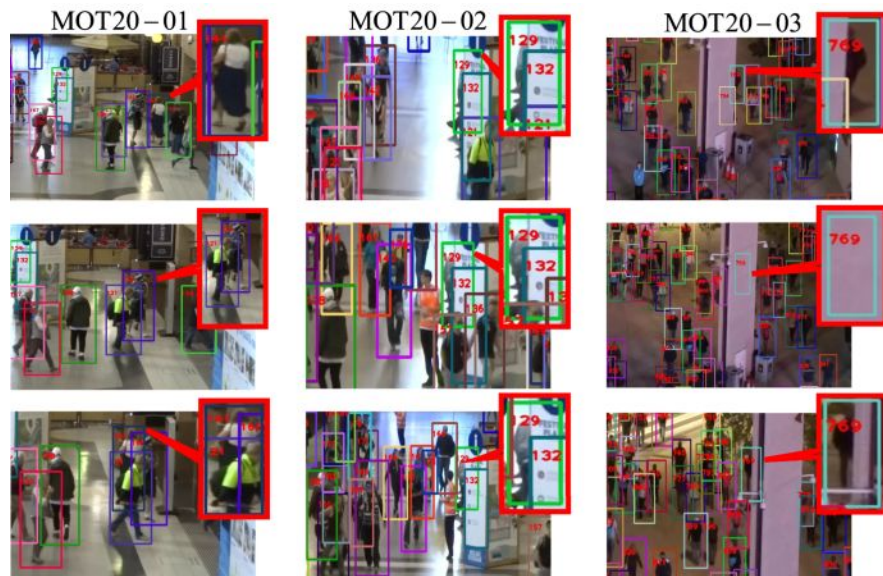
- YOLOX
 - Repository:
<https://github.com/Megvii-BaseDetection/YOLOX/blob/main/README.md>
 - Downloaded YOLOX-s weights
 - Best for real-time performance
- ByteTrack
 - Reminder: Multi-Object Tracking Algorithm, uses high and low scoring detection confidences
 - Repository:
<https://github.com/mikel-brostrom/boxmot>
- MOT20 Dataset
 - Full dataset download:
<https://motchallenge.net/data/MOT20/>
 - Utilized training data (4 videos): MOT20-01, MOT20-02, MOT20-03, MOT20-05

MOT20 Training Set

Sample	Name	FPS	Resolution	Length	Tracks	Boxes
	MOT20-05	25	1654x1080	3315 (02:13)	1211	751330
	MOT20-03	25	1173x880	2405 (01:36)	735	356728
	MOT20-02	25	1920x1080	2782 (01:51)	296	202215
	MOT20-01	25	1920x1080	429 (00:17)	90	26647
Total				8931 frm. (357 s.)	2332	1336920

Parallelization

- Frames must be processed in order
 - Cannot process two frames at once
- Detection must happen before tracking
 - Detect all frames (YOLOX)
 - Then track all detections (ByteTrack)
- Pixels in each frame can be processed in parallel
 - Detection
- Each detection box can be processed in parallel
 - Tracking



Targeted Accelerator

- NVIDIA Xavier NX
 - Symmetric CPU: 6 Camel Cores
 - GPU: 384 CUDA Cores, 48 Tensor Cores
Volta Generation
- NVIDIA Orin AGX
 - Symmetric CPU: 12 Cortex-A78AE
 - GPU: 2048 CUDA Cores, 64 Tensor Cores
Ampere Generation
- Preliminary Notes:
 - Xavier: **Fewer CUDA cores**, lower power, **less expensive** than Orin AGX
 - Orin AGX: **More CUDA cores**, higher power, **more expensive** than Xavier NX (*foreshadowing...*)

Xavier NX Functional Block Diagram

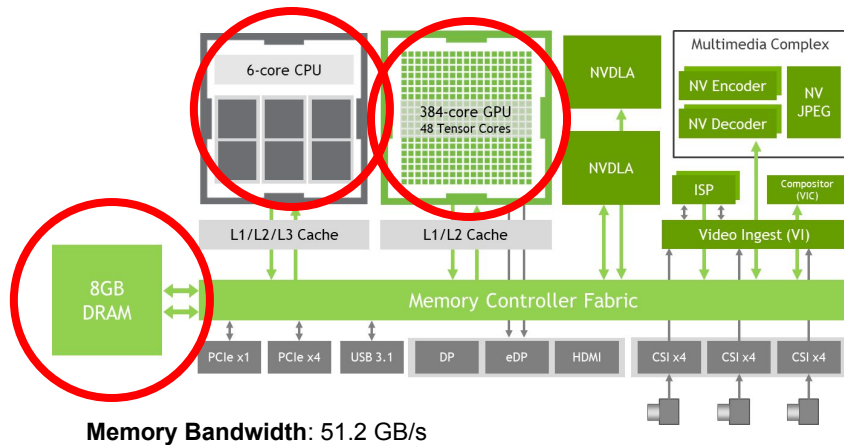
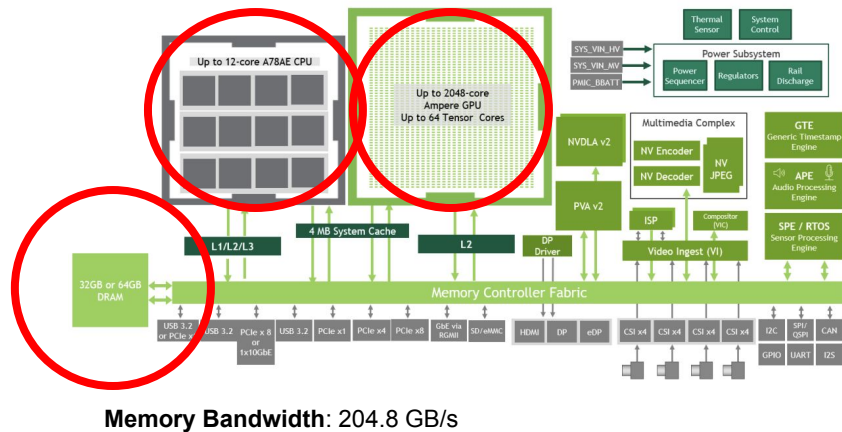


Figure 9: Jetson AGX Orin Series Functional Block Diagram

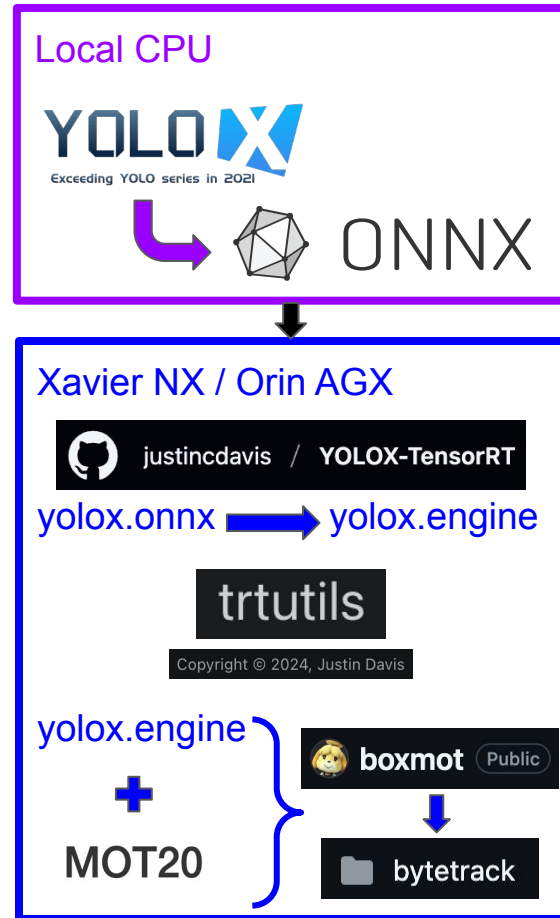


Acceleration Framework

- TensorRT / ONNX
 - Exported PyTorch model as ONNX
 - TensorRT used ONNX model to construct YOLOX model for NVIDIA devices
- CUDA
 - Both Xavier NX and Orin AGX: `pip install cuda-python==11.*`
 - Creating YOLOX model:
 - # initialize YOLOX model with CUDA
`yolo = YOLOX("yolox.engine", preprocessor="cuda")`
 - # preprocess the image
`tensor, ratio, padding = yolo.preprocess(img, method="cuda", no_copy=True)`
 - **TL;DR: Image processing done on GPU with CUDA, Object detection done on GPU with CUDA**

Overall Pipeline


- On local CPU:
 - Clone YOLOX
 - Export YOLOX model to ONNX format for TensorRT usage
 - scp yolox.onnx to Xavier NX
- SSH'ed to Xavier NX:
 - Clone YOLOX-TensorRT
 - Install trtutils (for using TensorRT to run YOLOX model)
 - Export yolox.onnx to TensorRT engine: yolox.engine
 - YOLOX inference (detection) on MOT20 videos using TensorRT engine with CUDA-accelerated preprocessing and tracking
 - Create Python v3.9 virtual environment
 - Install boxmot (tracking package)
 - Create ByteTrack object and feed in YOLOX outputted detections
- Orin AGX:
 - Nearly same process, needed a different yolox.engine (specific to device)



Challenges with Frameworks

- torch and torchvision
 - Jetson devices use custom firmware and drivers
 - Pre-built PyTorch and TorchVision versions (e.g., from pip) often don't match Jetson's CUDA version
 - We did PyTorch → ONNX → TensorRT for running on GPU
 - TensorRT 2x-10x faster than PyTorch on Jetson hardware
 - In our case, for running trained neural network model (YOLOX)
- NVRTC (NVIDIA Runtime Compilation)
 - CUDA library for compiling CUDA C++ code at runtime, opposed to build-time (runtime better for flexibility, performance)
 - Issues with version dependencies
 - RuntimeError: Failed to dlopen libnVRTC.so.12
 - CUDA-python v12 → **cuda-python==11.***

Running YOLOX model:



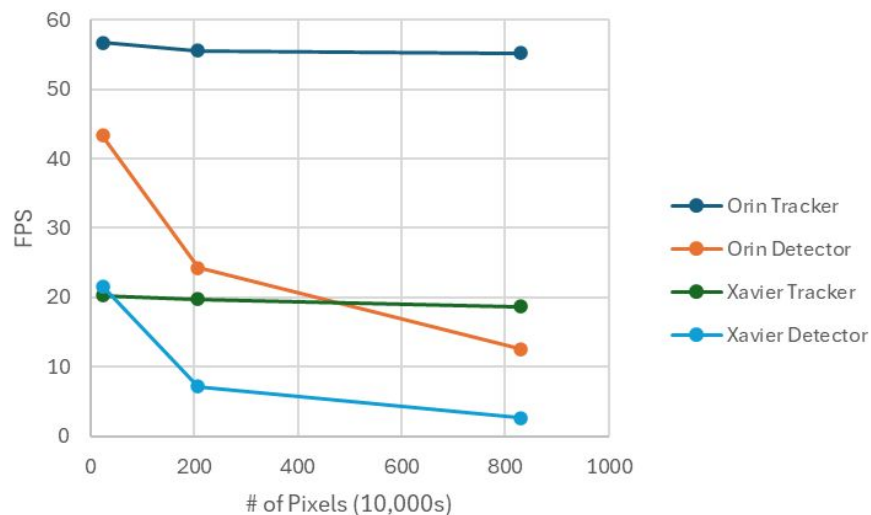
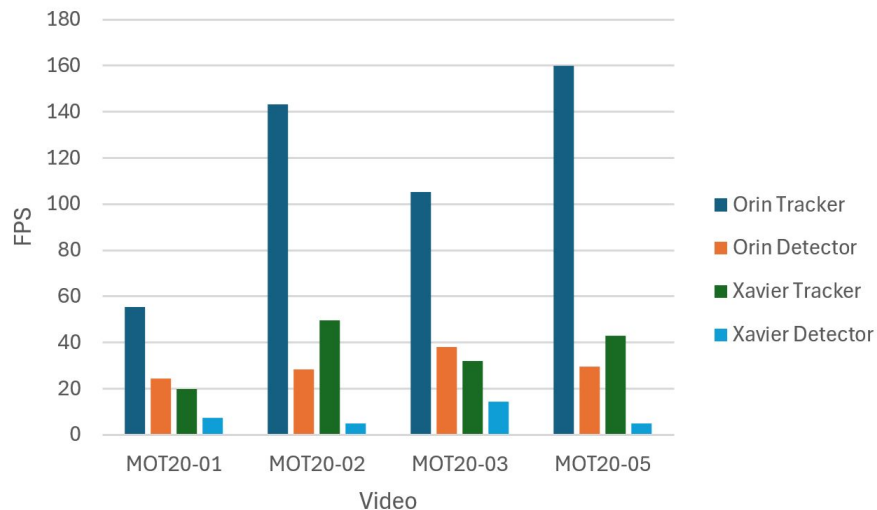
Performance Metrics

- Time/FPS
- Memory Usage
- Utilization of CPU and GPU
- CPU temperatures

Across different resolutions and videos

Measured separately for detector and tracker


Framerate



- Tracker is much faster
- Tracker is robust to resolution
- Tracker and detector heavily affected by complexity of scene

Memory Usage


- Orin uses more memory
 - It has more memory to use
- Detector uses (slightly) more memory
- No notable differences for different videos/resolutions



*Note: This is max RAM, but it is very comparable to average RAM



Memory Over Time

- Memory stays very steady after initial set up



CPU Usage


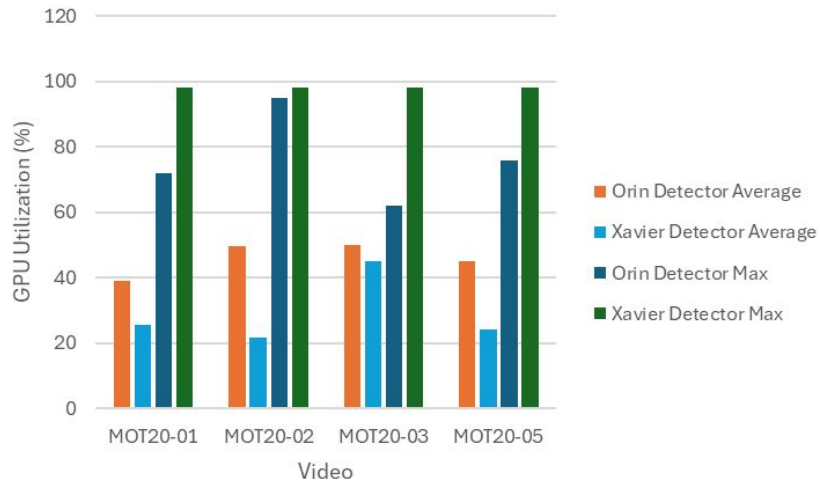
- Does not vary much based on input
- Much higher on Xavier (because Xavier's CPU capacity is lower)



GPU Usage


- More pixels decreases average utilization
- Lower average usage on Xavier (with a weaker GPU)

Both of these findings seem counter-intuitive, could be because CPU actions take up more time



Processor Utilization Over Time

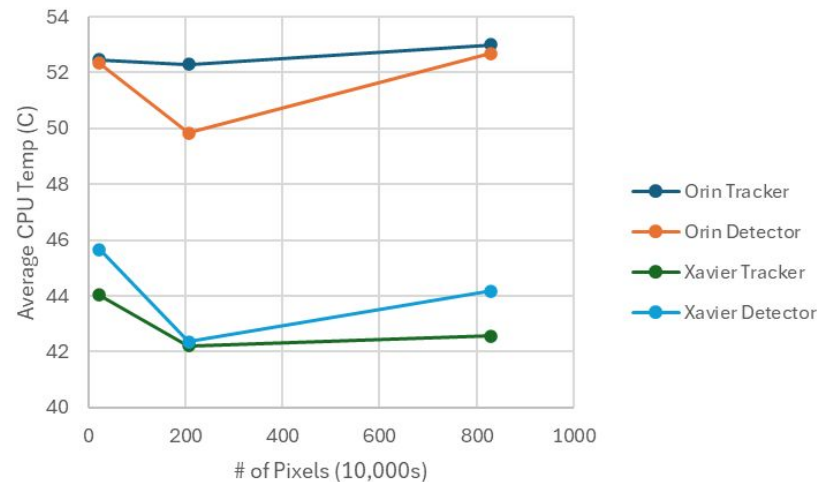
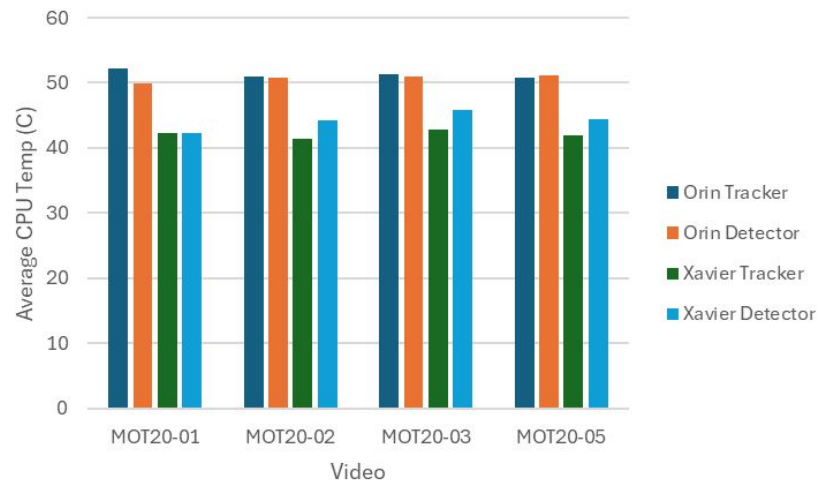
- GPUs very sporadic
- CPUs most utilized at start of tracking or detecting
 - Setup tasks



Temperature


- Orin runs hotter
 - More powerful CPU
- Seems to be an optimal resolution around 2,000,000 pixels (1080p)

*Note: ~2°C difference between average and max



Temperature Over Time

- Very steady over time
- No significant findings here



Future Directions and Takeaways

- CPU parts may be a bottleneck in detector
 - GPU has a lot of idle time
 - Future work could aim to improve parallelization if possible
- Detector takes up the most time (almost 2x tracker time)
 - Future work could look into faster detectors
 - Initial detection could use slower, more accurate detector. Then subsequent detections could use faster, less accurate detectors
- Future work could investigate domain-specific hardware
 - Special hardware could be developed/used for detections, tracker could still use CPU